

A DISTRIBUTED OPERATING SYSTEM FOR
PERMISSIONED BLOCKCHAINS

HYPERLEDGER FABRIC

BYZANTINE GENERALS' PROBLEM

- ▶ An agreement problem
- ▶ A group of generals wants to attack a city
- ▶ They must to know when to attack or when to retreat
- ▶ They cannot see or hear each other
- ▶ What do they need? → A PLAN
- ▶ We should build a network (nodes, communication channels, protocols)

BYZANTINE GENERALS' PROBLEM

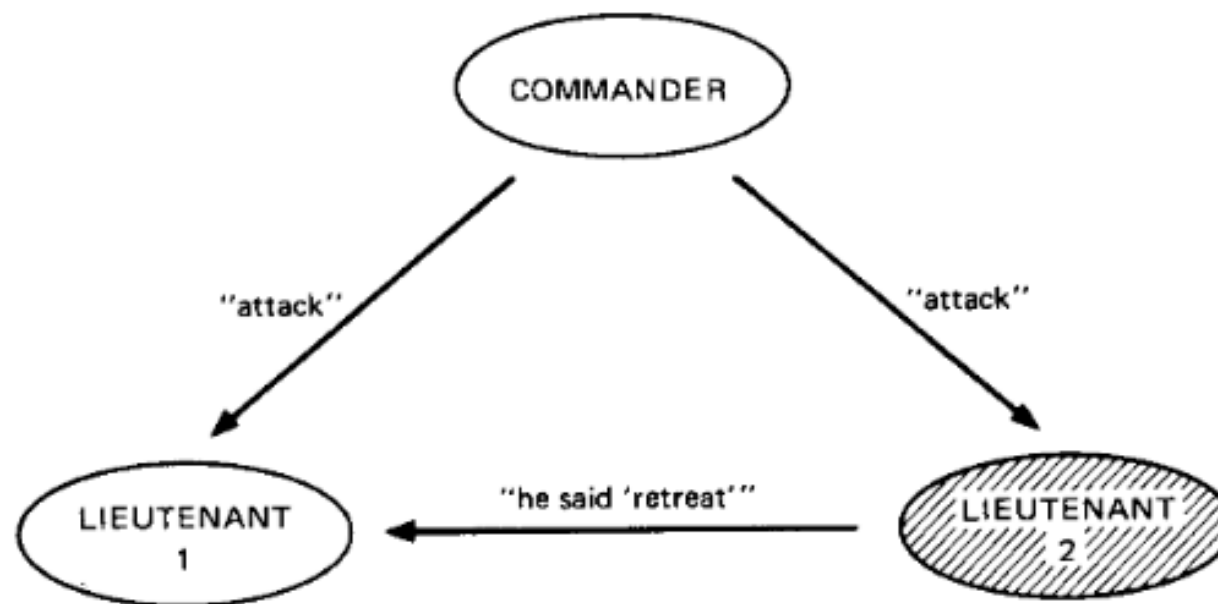


Fig. 1. Lieutenant 2 a traitor.

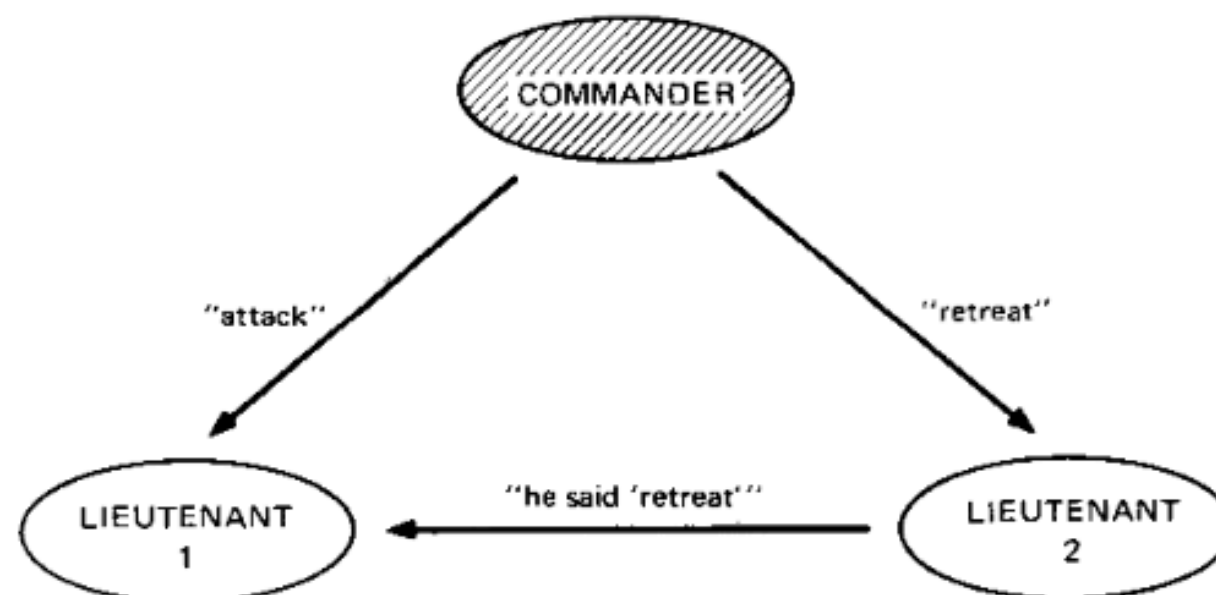


Fig. 2. The commander a traitor.

... take me to your commander

BYZANTINE GENERALS' PROBLEM IN PRACTICE

- ▶ Aircraft systems (Boeing 777 and 787 flight control systems)
- ▶ Spacecrafts (SpaceX Dragon flight system)
- ▶ Bitcoin

BLOCKCHAINS

- ▶ Simply, a blockchain is a huge ledger for recording transactions
- ▶ It is maintained within a distributed network of mutually untrusting peers (the generals)
- ▶ Every peer maintains a **copy** of the ledger
- ▶ They validate and order the transactions through a **consensus** protocol
- ▶ Blockchains have emerged with Bitcoin

BLOCKCHAINS

- ▶ In a public (*permissionless*) anyone can participate without a specific identity
- ▶ *Permissioned* blockchains run a blockchain (of course) among a set of *known, identified* participants. This is a way to secure the interactions among a group of entities that have a common goal but which do not fully trust each other
- ▶ Blockchains may execute arbitrary, programmable transaction logic in the form of **smart contracts** (as exemplified by Ethereum)

SMART CONTRACTS

- ▶ A smart contract functions as a trusted distributed application and gains its security from the blockchain and the underlying consensus among the peers
- ▶ Many existing smart contracts blockchains follow the blueprint of State-Machine Replication and implement so-called **active replication**: first, the transactions are ordered and propagated to all peers and second, each peer executes the transactions sequentially

SMART CONTRACTS

- ▶ Prior permissioned blockchains suffer from many limitations:
 - ▶ Smart contract must be written in a *fixed, non-standard, or domain-specific language*
 - ▶ The sequential execution of transaction by all peers *limits performance*
 - ▶ Transaction must be *deterministic*, which can be difficult to be ensured programmatically
 - ▶ Every smart contract runs on *all peers*, which is at odds with *confidentiality*

OK, LET'S SEE FABRIC

- ▶ Fabric introduces a new blockchain architecture aiming at *flexibility, scalability and confidentiality*
- ▶ Fabric support the execution of distributed applications written in *standard programming languages*
- ▶ Fabric is the first *distributed operating system* for permissioned blockchains

FABRIC

- ▶ The architecture of Fabric follows the *execute-order-validate* paradigm



FABRIC

- ▶ This design departs radically from the *order-execute* paradigm
- ▶ It combines the two approaches to replication, *passive* and *active*:
 - ▶ Every transaction is executed (endorsed) only by a subset of peers, which allows for parallel execution (passive)
 - ▶ The transaction's effects on the ledger state are only written after reaching consensus of a total order among them (active)

FABRIC

- ▶ This *hybrid replication design*, which mixes passive and active replication in the Byzantine model, and the *execute-order-validate* paradigm, represent the most innovation in Fabric architecture

ORDER-EXECUTE ARCHITECTURE

- ▶ All previous blockchain systems follows *order-execute architecture*
- ▶ Let's take Ethereum for example:
 1. every peer assembles a block containing valid transactions
 2. the peer tries to solve the puzzle
 3. if the peer is *lucky* and solves the puzzle it disseminates the block to the network
 4. every peer receiving the block validates the solution *and* all transaction in the block
- ▶ Simply said, every peer **repeats** the execution of the lucky peer from its first step
- ▶ If this is not enough, all transactions must be executed **sequentially**

LIMITATIONS OF ORDER-EXECUTE ARCHITECTURE

- ▶ Sequentially execution
 - ▶ It limits the effective throughput that can be achieved
 - ▶ Since the throughput is inversely proportional to the execution latency, this may become a performance bottleneck
 - ▶ A *Denial-of-Service* attack could simply introduce smart contracts that take very long time to execute

LIMITATIONS OF ORDER-EXECUTE ARCHITECTURE

- ▶ Non-deterministic code
 - ▶ This is usually addressed by programming blockchains in domain-specific languages (e.g. Ethereum Solidity) that are expressive enough for their applications but limited for deterministic execution
 - ▶ Only one non-deterministic contract created with malicious intent is enough to bring the entire blockchain to a halt

LIMITATIONS OF ORDER-EXECUTE ARCHITECTURE

- ▶ Confidentiality of execution
 - ▶ Many permissioned systems run all smart contract on all peers
 - ▶ However, many intended use cases for permissioned blockchains require confidentiality